

One-line answers, armed. Lead with the bold line; the green bit is the mic-drop.

TECHNICAL — "HOW DO YOU KNOW?"

Q How do you know my quota?

It's an **estimate**, not a live link. Anchored at onboarding from the BUDI95 app/receipt, projected forward by logged fills. **Missed a fill? One tap "Correct balance" re-anchors** — so drift never poisons the verdict.

Q What happens after the quota's gone?

Same engine, second clock. The **Float Regime** bridges to the next known weekly price post (Wednesday). **We surface the known reprice date — we never predict the price.**

Q Isn't the AI just formatting a calculator?

Code computes & validates every number; the AI only voices it — it literally can't show a number the engine didn't produce. The non-calculator part is **Ask JimatJalan**: free-form reasoning over your real state ("drive to JB?" → **RM85.96**, grounded).

Q Where's the prediction the brief wants?

We predict your spend, not the price. RM119 next month from your burn × price. **Predicting the price would be a fragile lie; predicting your spend is honest arithmetic.**

Q How is "save RM58" real?

vs filling the **whole tank at today's price**: 28 litres you'd buy at market now, bought subsidised after the reset instead, × the **RM2.08** subsidy gap. We show the subsidised+market split — fully auditable.

Q How do you know it's "hard acceleration"?

We don't assert it. Code picks from a **bounded cause menu** and labels them "likely culprits." **Never "we measured your driving" — inference, honestly framed.**

Q Accessibility?

Visible focus rings, AA contrast, ≥48px tap targets, colour never the only signal, reduced-motion respected, semantic headings — and it's an **installable PWA**.

Q What if the AI / wifi dies on stage?

It can't go blank. Every surface falls back to a deterministic, number-exact voice. **The decision is code; only the wording needs the model.**

BUSINESS — "WHY YOU, WHY NOW, WHY IT LASTS"

Q What's the moat? Anyone can build this.

Not the formula — the **segment-tuned model** (Quota Cap 200/800, bridge-vs-bank) generic wallets don't have, plus **per-driver data** (baseline + fill cadence → switching cost) feeding a **B2B fleet** play. And the can't-lie guard is a trust edge a wrapper can't claim.

Q Won't Setel / TnG / Maybank just bundle it free?

Distribution beats features only on the **mass consumer**. We win the **high-pain e-hailing/gig wedge + fleet B2B** — segments their generic wallets don't model. Land the wedge, sell the fleet.

Q BUDI95 could change or vanish — that's your whole premise.

It already **changed twice in 2026**. The two-clock engine is **config, not premise** — a policy change is an edit, not a rewrite. **We're a fuel-spend optimiser for any subsidy/price regime**; BUDI95 is just today's instance.

Q Revenue & financial sustainability?

Freemium (premium analytics) · **B2B fleet** (per-driver/mo + sustainability reporting) · **partnerships** (stations, insurers). **Marginal cost ≈ a fraction-of-a-sen LLM call; most value served at zero AI cost.**

Q Market size & demand?

National RON95 base; the wedge is the large, high-spend, high-churn **e-hailing/gig** segment. **Why now: the BUDI95 quota model is brand-new and volatile** — quota literacy suddenly matters. *(Drop in verified driver/spend figures.)*

Q Scalability?

Config-driven engine (new segment = data, not code), swappable Store adapter (→ Postgres), cheap unit economics, and it degrades to a zero-cost deterministic mode under load.

If cornered, fall back to three lines: ① It's an estimate, re-anchorable. ② Code owns the numbers, the AI owns the voice — it can't lie. ③ Config, not premise — it survives any policy.